

Iteration 7: **PULL** (i.e., arm)

8 arms in pool | 60 unseen left

best dev 56.9 → 58.5%


while several high-activity known arms **remain genuinely active and broad** ... the known-arm set now contains **multiple severe, plausibly fixable** population-level weaknesses ... another DRAW would likely defer repairs that are **already well diagnosed**.

Iteration 2: **DRAW** (i.e., unseen)

2 arms in pool | 72 unseen left

best dev 55.4 → 56.9%

Only two known arms exist and both come from the same scenario ... the **failure surface is poorly mapped**, so exploring unseen scenarios is **more valuable** for this iteration.



Iteration 30: **PULL** (i.e., arm)

24 arms in pool | 33 unseen left

best dev 61.5 → 63.1%

the latest evaluation exposed a **fresh, severe, trace-backed known arm** ... older known arms are inactive or already resolved, and the trace-less timeout arm is severe but much less diagnosable ... the current known arm provides a **clearer expected learning opportunity**.

Iteration 24: **DRAW** (i.e., unseen)

16 arms in pool | 45 unseen left

best dev 58.5 → 61.5%

the proven C7/C22 executor fixes **already cover most previously pulled arms** ... arms are either fixed, **recently all-pass**, or **repeatedly inactive** on later candidates ... **discovery value is higher** than spending this iteration on a **mostly exhausted known-arm set**.